

## Explanation of the below CNN code block:

```
n_classes = 24
kernel_size = 3
flattened_img_size = 75 * 3 * 3
model = nn.Sequential(
    # First convolution
    nn.Conv2d(IMG_CHS, 25, kernel_size, stride=1, padding=1), # 25 x 28 x 28
    nn.BatchNorm2d(25),
    nn.ReLU(),
    nn.MaxPool2d(2, stride=2), # 25 x 14 x 14
    # Second convolution
    nn.Conv2d(25, 50, kernel_size, stride=1, padding=1), # 50 x 14 x 14
    nn.BatchNorm2d(50),
    nn.ReLU(),
    nn.Dropout(.2),
    nn.MaxPool2d(2, stride=2), # 50 x 7 x 7
    # Third convolution
    nn.Conv2d(50, 75, kernel_size, stride=1, padding=1), # 75 x 7 x 7
    nn.BatchNorm2d(75),
    nn.ReLU(),
    nn.MaxPool2d(2, stride=2), # 75 x 3 x 3
    # Flatten to Dense
    nn.Flatten(),
    nn.Linear(flattened_img_size, 512),
    nn.Dropout(.3),
    nn.ReLU(),
    nn.Linear(512, n_classes)
)
```

For each convolution and pooling layer, the output shape depends on the input shape, kernel size, stride, and padding.

### 1. Convolution Layer Output Shape

The output shape for a 2D convolutional layer can be calculated as follows for each dimension (height and width):

$$\text{Output Dimension} = \frac{\text{Input Dimension} - \text{Kernel Size} + 2 \times \text{Padding}}{\text{Stride}} + 1$$

where:

- **Input Dimension** is the height or width of the input.
- **Kernel Size** is the size of the convolutional kernel (filter).
- **Padding** is the number of pixels added around the input.
- **Stride** is the step size of the kernel.

The number of output channels (depth) is determined by the number of filters applied, not affected by the height and width calculation.

## 2. Max Pooling Layer Output Shape

For max pooling layers, the output shape can be calculated similarly:

$$\text{Output Dimension} = \frac{\text{Input Dimension} - \text{Pool Size}}{\text{Stride}} + 1$$

where **Pool Size** is the size of the pooling window.

## Walkthrough of Each Layer

Let's go layer by layer using these formulas.

---

### Layer 1: First Convolution Block

1. **Input Shape:** `IMG_CHS x 28 x 28`
  2. **Conv2d Layer:** `nn.Conv2d(IMG_CHS, 25, kernel_size=3, stride=1, padding=1)`
    - Kernel Size = 3
    - Padding = 1
    - Stride = 1
    - **Output Shape Calculation:**
      - Height and Width:  $\frac{28-3+2 \times 1}{1} + 1 = 28$
      - The output shape is `25 x 28 x 28` (25 is the number of output channels).
  3. **MaxPool2d Layer:** `nn.MaxPool2d(2, stride=2)`
    - Pool Size = 2
    - Stride = 2
    - **Output Shape Calculation:**
      - Height and Width:  $\frac{28-2}{2} + 1 = 14$
      - The output shape is `25 x 14 x 14`.
- 

### Layer 2: Second Convolution Block

1. **Input Shape:** `25 x 14 x 14`
2. **Conv2d Layer:** `nn.Conv2d(25, 50, kernel_size=3, stride=1, padding=1)`
  - Kernel Size = 3
  - Padding = 1
  - Stride = 1

- **Output Shape Calculation:**
    - Height and Width:  $\frac{14-3+2 \times 1}{1} + 1 = 14$
    - The output shape is  $50 \times 14 \times 14$ .
  - 3. **MaxPool2d Layer:** `nn.MaxPool2d(2, stride=2)`
    - Pool Size = 2
    - Stride = 2
    - **Output Shape Calculation:**
      - Height and Width:  $\frac{14-2}{2} + 1 = 7$
      - The output shape is  $50 \times 7 \times 7$ .
- 

### Layer 3: Third Convolution Block

1. **Input Shape:**  $50 \times 7 \times 7$
  2. **Conv2d Layer:** `nn.Conv2d(50, 75, kernel_size=3, stride=1, padding=1)`
    - Kernel Size = 3
    - Padding = 1
    - Stride = 1
    - **Output Shape Calculation:**
      - Height and Width:  $\frac{7-3+2 \times 1}{1} + 1 = 7$
      - The output shape is  $75 \times 7 \times 7$ .
  3. **MaxPool2d Layer:** `nn.MaxPool2d(2, stride=2)`
    - Pool Size = 2
    - Stride = 2
    - **Output Shape Calculation:**
      - Height and Width:  $\frac{7-2}{2} + 1 = 3$
      - The output shape is  $75 \times 3 \times 3$ .
- 

### Flatten Layer

1. **Input Shape:**  $75 \times 3 \times 3$
  2. **Flattening:** Converts this 3D tensor to a 1D tensor with size  $75 * 3 * 3 = 675$ .
- 

### Fully Connected (Dense) Layers

1. **Linear Layer:** `nn.Linear(675, 512)`
  - Input: 675 (flattened image size)

- Output: 512 (size of the hidden layer)
- 2. **Final Linear Layer:** `nn.Linear(512, n_classes)`
  - Input: 512
  - Output: `n_classes` (24)